

Chapter 3. Random Variables and Distributions

Basic idea of introducing random variables: represent outcomes and/or random events by numbers.

3.1 Random variables and Distributions.

Definition. A random variable is a function defined on the sample space Ω , which assigns a real number $X(\omega)$ to each outcome $\omega \in \Omega$.

Example 1. Flip a coin 10 times. We may define random variables (r.v.s) as follows:

X_1 = no. of heads,

X_2 = no. of flips required to have the first head,

X_3 = no. of 'HT'-pairs,
 X_4 = no. of tails.

For $\omega = HHTHHTHHTT$, $X_1(\omega) = 6$, $X_2(\omega) = 1$, $X_3(\omega) = 3$, $X_4(\omega) = 4$.
Note $X_1 \equiv 10 - X_4$.

Remark. The values of a r.v. varies and cannot be pre-determined before an outcome occurs.

Definition. For any r.v. X , its (cumulative) distribution function (CDF) is defined as $F_X(x) = P(X \leq x)$.

Example 2. Toss a fair coin twice and let X be the number of heads. Then

$$P(X = 0) = P(X = 2) = 1/4, \quad P(X = 1) = 1/2.$$

Hence its CDF is $F_X(x) = \begin{cases} 0 & x < 0, \\ 1/4 & 0 \leq x < 1, \\ 3/4 & 1 \leq x < 2, \\ 1 & x \geq 2. \end{cases}$

Note. (i) $F_X(x)$ is right continuous, non-decreasing, and defined for all $x \in (-\infty, \infty)$. For example, $F_X(1.1) = 0.75$.

(ii) The CDF is a non-random function.

(iii) If $F(\cdot)$ is the CDF of r.v. X , we simply write $X \sim F$.

Properties of CDF. A function $F(\cdot)$ is a CDF iff

- (i) F is non-decreasing: $x_1 < x_2$ implies $F(x_1) \leq F(x_2)$,
- (ii) F is normalized: $\lim_{x \rightarrow -\infty} F(x) = 0$, $\lim_{x \rightarrow \infty} F(x) = 1$,
- (iii) F is right continuous: $\lim_{y \downarrow x} F(y) = F(x)$.

Probabilities from CDF

$$(a) P(X > x) = 1 - F(x)$$

$$(b) P(x < X \leq y) = F(y) - F(x)$$

$$(c) P(X < x) = \lim_{h \downarrow 0} F(x - h) \equiv F(x-)$$

$$(d) P(X = x) = F(x) - F(x-).$$

Note. It is helpful for understanding (c) & (b) to revisit Example 2.

3.2 Discrete random variables

If r.v. X only takes some isolated values, X is called a discrete r.v. Its CDF is called a discrete distribution.

Definition. For a discrete r.v. X taking values $\{x_1, x_2, \dots\}$, we define the probability function (or probability mass function) as

$$f_X(x_i) = P(X = x_i), \quad i = 1, 2, \dots$$

Obviously, $f_X(x_i) \geq 0$ and $\sum_i f_X(x_i) = 1$.

It is often more convenient to list a probability function in a table:

X	x_1	x_2	$\dots\dots$
Probability	$f_X(x_1)$	$f_X(x_2)$	$\dots\dots$

Example 2 (continue). The probability function is be tabulated:

X	0	1	2
Probability	$1/4$	$1/2$	$1/4$

Expectation or Mean EX or $E(X)$: a measure for the ‘center’, ‘average value’ of a r.v. X , and is often denoted by μ .

For a discrete r.v. X with probability function $f_X(x)$,

$$\mu = EX = \sum_i x_i f_X(x_i).$$

Variance $\text{Var}(X)$: a measure for variation, uncertainty or ‘risk’ of a r.v. X , is often denoted by σ^2 , while σ is called **standard deviation** of X .

For a discrete r.v. X with probability function $f_X(x)$,

$$\sigma^2 = \text{Var}(X) = \sum_i (x_i - \mu)^2 f_X(x_i) = \sum_i x_i^2 f_X(x_i) - \mu^2.$$

The k -th moment of X : $\mu_k \equiv E(X^k) = \sum_i x_i^k f_X(x_i)$, $k = 1, 2, \dots$.

Obviously, $\mu = \mu_1$, and $\sigma^2 = \mu_2 - \mu_1^2$.

Some important discrete distributions

Convention. We often use upper case letters X, Y, Z, \dots to denote r.v.s, and lower case letters x, y, z, \dots to denote the values of r.v.s. In contrast letters a, b or A, B are often used to denote (non-random) constants.

Degenerate distribution: $X \equiv a$, i.e. $F_X(x) = 1$ for any $x \geq a$, and 0 otherwise.

It is easy to see that $\mu = a$ and $\sigma^2 = 0$.

Bernoulli distribution: X is binary, $P(X = 1) = p$ and $P(X = 0) = 1 - p$, where $p \in [0, 1]$ is a constant. It represents the outcome of flipping a coin.

$$\mu = 1 \cdot p + 0 \cdot (1 - p) = p, \quad \sigma^2 = p(1 - p).$$

Note. Bernoulli trial refers to an experiment of flipping a coin repeatedly.

Binomial distribution $\text{Bin}(n, p)$: X takes values $0, 1, \dots, n$ only with the probability function

$$f_X(x) = \frac{n!}{x!(n-x)!} p^x (1-p)^{n-x}, \quad x = 0, 1, \dots, n.$$

Theorem. If we toss a coin n times, let X be the number of heads. Then $X \sim \text{Bin}(n, p)$, where p is the probability that head occurs in tossing the coin once.

Proof. Let $\omega = HT HHT \dots H$ denote an outcome of n tosses. Then $X = k$ iff there are k 'H' and $(n - k)$ 'T' in ω . Therefore the probability of such a ω is $p^k (1 - p)^{n-k}$. Since those k H's may occur in any n positions of the sequence, there are $\binom{n}{k}$ such ω 's. Hence

$$P(X = k) = \binom{n}{k} p^k (1 - p)^{n-k} = \frac{n!}{k!(n-k)!} p^k (1 - p)^{n-k}, \quad k = 0, 1, \dots, n.$$

Let us check if the probability function above is well defined. Obviously $P(X = k) \geq 0$, furthermore

$$\sum_{k=0}^n P(X = k) = \sum_{k=0}^n \binom{n}{k} p^k (1-p)^{n-k} = \{p + (1-p)\}^n = 1^n = 1.$$

Let us work out the mean and the variance for $X \sim \text{Bin}(n, 1-p)$.

$$\begin{aligned} \mu &= \sum_{k=0}^n k \frac{n!}{k!(n-k)!} p^k (1-p)^{n-k} = \sum_{k=1}^n k \frac{n!}{k!(n-k)!} p^k (1-p)^{n-k} \\ &= \sum_{j=0}^{n-1} np \frac{(n-1)!}{j!(n-1-j)!} p^j (1-p)^{n-1-j} \\ &= np \sum_{j=0}^m \frac{m!}{j!(m-j)!} p^j (1-p)^{m-j} = np. \end{aligned}$$

Note that $\sigma^2 = E(X^2) - \mu^2 = E\{X(X-1)\} + \mu - \mu^2$. We need to work out

$$\begin{aligned} E\{X(X-1)\} &= \sum_{k=0}^n k(k-1) \frac{n!}{k!(n-k)!} p^k (1-p)^{n-k} \\ &= \sum_{k=2}^n n(n-1) p^2 \frac{(n-2)!}{(k-2)! \{(n-2)-(k-2)\}} p^{k-2} (1-p)^{\{(n-2)-(k-2)\}} \\ &= n(n-1) p^2 \sum_{j=0}^{n-2} \frac{(n-2)!}{j! \{(n-2)-j\}} p^j (1-p)^{\{(n-2)-j\}} = n(n-1) p^2. \end{aligned}$$

This gives $\sigma^2 = n(n-1)p^2 + np - (np)^2 = np(1-p)$.

By the above theorem, we can see immediately

- (i) If $X \sim \text{Bin}(n, p)$, $n - X \sim \text{Bin}(n, 1 - p)$.
- (ii) If $X \sim \text{Bin}(n, p)$, $Y \sim \text{Bin}(m, p)$, and X and Y are independent, then $X + Y \sim \text{Bin}(n + m, p)$.

Furthermore, let $Y_i = 1$ if the i -th toss yields H, and 0 otherwise. Then Y_1, \dots, Y_n are *independent* Bernoulli r.v.s with mean p and variance $p(1 - p)$. Since $X = Y_1 + \dots + Y_n$, we notice

$$EX = \sum_{i=1}^n EY_i = np, \quad \text{Var}(X) = \sum_{i=1}^n \text{Var}(Y_i) = np(1 - p).$$

This is a much easier way to derived the means and variances for binomial distributions, which is based on the following general properties.

(i) For any r.v.s ξ_1, \dots, ξ_n , and any constants a_1, \dots, a_n ,

$$E\left(\sum_{i=1}^n a_i \xi_i\right) = \sum_{i=1}^n a_i E(\xi_i).$$

(ii) If, in addition, ξ_1, \dots, ξ_n are *independent*,

$$\text{Var}\left(\sum_{i=1}^n a_i \xi_i\right) = \sum_{i=1}^n a_i^2 \text{Var}(\xi_i).$$

Independence of random variables. The r.v.s ξ_1, \dots, ξ_n are independent if

$$P(\xi_1 \leq x_1, \dots, \xi_n \leq x_n) = P(\xi_1 \leq x_1) \times \dots \times P(\xi_n \leq x_n)$$

for any x_1, \dots, x_n .

Moment generate function (MGF) of r.v. X :

$$\psi_X(t) = E(e^{tX}), \quad t \in (-\infty, \infty).$$

(i) It is easy to see that $\psi'_X(0) = E(X) = \mu$. In general $\psi_X^{(k)}(0) = E(X^k) = \mu_k$.

(ii) If $Y = a + bX$, $\psi_Y(t) = E(e^{(a+bX)t}) = e^{at}\psi_X(bt)$.

(iii) If X_1, \dots, X_n are independent, $\psi_{\sum_i X_i}(t) = \prod_{i=1}^n \psi_{X_i}(t)$, and vice versa

If X is discrete, $\psi_X(t) = \sum_i e^{x_i t} f_X(x_i)$.

To generate a r.v. from $\text{Bin}(n, p)$, we can flip a coin (with p -probability for H) n times, and count the number of heads. However R can do the flipping for us much more efficiently:

```
> rbinom(10, 100, 0.1) # generate 10 random numbers from \Bin(100, 0.1)
[1]  8 11  9  9 18  7  5  5  3  7
> rbinom(10, 100, 0.1) # do it again, obtain different numbers
[1] 11 13  6  7 11  9  9  9 12 10
> x <- rbinom(10, 100, 0.7); x; mean(x)
[1] 66 77 67 66 64 68 70 68 72 72
[1] 69 # mean close to np=70
> x <- rbinom(10, 100, 0.7); x; mean(x)
[1] 70 73 72 70 68 69 70 66 79 71
[1] 70.8
```


Note that `rbinom(10000, 1, 0.5)` is equivalent to toss a fair coin 10000 times:

```
> y <- rbinom(10000, 1, 0.5); length(y); table(y)
[1] 10000
y
 0      1
4990 5010 # about a half times with head
```

You may try with smaller sample size, such as

```
> y <- rbinom(10, 1, 0.5); length(y); table(y)
[1] 10
y
0 1
3 7 # 7 heads and 3 tails
```

Also try out pbinom (CDF), dbinom (probability function), qbinom (quantile) for Binomial distributions.

Geometric Distribution Geom(p): X takes all positive integer values with probability function

$$P(X = k) = (1 - p)^{k-1} p, \quad k = 1, 2, \dots$$

Obviously, X is the number of tosses required in a Bernoulli trial to obtain the first head.

$$\mu = \sum_{k=1}^{\infty} k(1 - p)^{k-1} p = -p \frac{d}{dp} \sum_{k=1}^{\infty} (1 - p)^k = -p \frac{d}{dp} (1/p) = 1/p,$$

and it can be shown that $\sigma^2 = (1 - p)/p^2$.

Using the MGF provides an alternative way to find mean and variance: for

$t < -\log(1 - p)$ (i.e. $e^t(1 - p) < 1$),

$$\begin{aligned}\psi_X(t) &= E(e^{tX}) = \sum_{i=1}^{\infty} e^{ti}(1 - p)^{i-1}p = \frac{p}{1 - p} \sum_{i=1}^{\infty} \{e^t(1 - p)\}^i \\ &= \frac{p}{1 - p} \frac{e^t(1 - p)}{1 - e^t(1 - p)} = \frac{pe^t}{1 - e^t(1 - p)} = \frac{p}{e^{-t} - 1 + p}.\end{aligned}$$

Now $\mu = \psi'_X(0) = \left[\frac{pe^{-t}}{(e^{-t} - 1 + p)^2} \right]_{t=0} = 1/p$, and

$$\mu_2 = \psi''_X(0) = \left[\frac{2pe^{-2t}}{(e^{-t} - 1 + p)^3} - \frac{pe^{-t}}{(e^{-t} - 1 + p)^2} \right]_{t=0} = 2/p^2 - 1/p.$$

Hence $\sigma^2 = \mu_2 - \mu^2 = (1 - p)/p^2$.

The R functions for $\text{Geom}(p)$: `rgeom`, `dgeom`, `pgeom` and `qgeom`.

Poisson Distribution $\text{Poisson}(\lambda)$: X takes all non-negative integers with probability function

$$P(X = k) = \frac{\lambda^k}{k!} e^{-\lambda}, \quad k = 0, 1, 2, \dots,$$

where $\lambda > 0$ is a constant, called parameter.

The MGF $X \sim \text{Poisson}(\lambda)$:

$$\psi_X(t) = \sum_{k=0}^{\infty} e^{kt} \frac{\lambda^k}{k!} e^{-\lambda} = e^{-\lambda} \sum_{k=0}^{\infty} \frac{(e^t \lambda)^k}{k!} = e^{-\lambda} e^{e^t \lambda} = \exp\{\lambda(e^t - 1)\}.$$

Hence

$$\mu = \psi'_X(0) = [\exp\{\lambda(e^t - 1)\} \lambda e^t]_{t=0} = \lambda,$$

$$\mu_2 = \psi''_X(0) = [\exp\{\lambda(e^t - 1)\} \lambda e^t + \exp\{\lambda(e^t - 1)\} (\lambda e^t)^2]_{t=0} = \lambda + \lambda^2.$$

Therefore $\sigma^2 = \mu_2 - \mu^2 = \lambda$.


Remark. For Poisson distributions, $\mu = \sigma^2$.

The R functions for Poisson(λ): `rpois`, `dpois`, `ppois` and `qpois`.

To understand the role of the parameter λ , we plot the probability function of Poisson(λ) for different values of λ .

```
> x <- c(0:20)
> plot(x,dpois(x,1),type='o',xlab='k', ylab='probability function')
> text(2.5,0.35, expression(lambda==1))
> lines(x,dpois(x,2),typ='o',lty=2, pch=2)
> text(3.5,0.25, expression(lambda==2))
> lines(x,dpois(x,8),typ='o',lty=3, pch=3)
> text(7.5,0.16, expression(lambda==8))
> lines(x,dpois(x,15),typ='o',lty=4, pch=4)
> text(14.5, 0.12, expression(lambda==15))
```

Plots of $\lambda^k e^{-\lambda} / k!$ against k



Three ways of computing probability and distribution functions:

- calculators — for simple calculation
- statistical tables — for, e.g. the final exam
- R — for serious tasks such as real application

3.2 Continuous random variables

A r.v. X is *continuous* if there exists a function $f_X(\cdot) \geq 0$ such that

$$P(a < X < b) = \int_a^b f_X(x) dx, \quad \forall a < b.$$

We $f_X(\cdot)$ the *probability density function* (PDF) or, simply, density function. Obviously

$$F_X(x) = \int_{-\infty}^x f_X(u) du.$$

Properties of continuous random variables

(i) $F_X(x) = P(X \leq x) = P(X < x)$, i.e. $P(X = x) = 0 \neq f_X(x)$.

(ii) The PDF $f_X(\cdot) \geq 0$, and $\int_{-\infty}^{\infty} f_X(x)dx = 1$.

(iii) $\mu = E(X) = \int_{-\infty}^{\infty} xf_X(x)dx$,

$$\sigma^2 = \text{Var}(X) = \int_{-\infty}^{\infty} (x - \mu)^2 f_X(x)dx = \int_{-\infty}^{\infty} x^2 f_X(x)dx - \mu^2.$$

Furthermore the MGF of X is equal to

$$\psi_X(t) = \int_{-\infty}^{\infty} e^{tx} f_X(x)dx.$$

Some important continuous distributions

Uniform distribution $U(a, b)$: X takes any values between a and b equally likely. Its PDF is

$$f(x) = \begin{cases} \frac{1}{b-a} & a \leq x \leq b \\ 0 & \text{otherwise.} \end{cases}$$

Then the CDF is

$$F(x) = \int_{-\infty}^x f(u)du = \begin{cases} 0 & x < a, \\ \frac{1}{b-a} \int_a^x du = \frac{x-a}{b-a} & a \leq x \leq b, \\ 1 & x > b. \end{cases},$$

and

$$\mu = \int_a^b \frac{x dx}{b-a} = \frac{a+b}{2}, \quad \mu_2 = \int_a^b \frac{x^2 dx}{b-a} = \frac{b^3 - a^3}{3(b-a)} = \frac{a^2 + ab + b^2}{3}$$

Hence $\sigma^2 = \mu_2 - \mu^2 = (b-a)^2/12$.

R-functions related to uniform distributions: runif, dunif, punif, qunif.

Quantile. For a given CDF $F(\cdot)$, its quantile function is defined as

$$F^{-1}(p) = \inf\{x : F(x) \geq p\}, \quad p \in [0, 1]$$

```
> x <- c(1, 2.5, 4)
> punif(x, 2, 3)      # CDF of U(2, 3) at 1, 2.5 and 4
[1] 0 0.5 1
> dunif(x, 2, 3)      # PDF of U(2, 3) at 1, 2.5 and 4
[1] 0 1 0
> qunif(0.5, 2, 3)    # quantile of U(2, 3) at p=0.5
[1] 2.5
```

Normal Distribution $N(\mu, \sigma^2)$: the PDF is

$$f(x) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{1}{2\sigma^2}(x - \mu)^2 \right\}, \quad -\infty < x < \infty,$$

where $\mu \in (-\infty, \infty)$ is the ‘centre’ (or mean) of the distribution, and $\sigma > 0$ is the ‘spread’ (standard deviation).

Remarks. (i) The most important distribution in statistics: Many phenomena in nature have approximately normal distributions. Furthermore, it provides asymptotic approximations for the distributions of sample means (Central Limit Theorem).

(ii) If $X \sim N(\mu, \sigma^2)$, $EX = \mu$, $\text{Var}(X) = \sigma^2$, and $\psi_X(t) = e^{\mu t + \sigma^2 t^2/2}$.

We compute $\psi_X(t)$ below, the idea is applicable in general.

$$\begin{aligned}
 \psi_X(t) &= \frac{1}{\sqrt{2\pi}\sigma} \int e^{tx} e^{-\frac{1}{2\sigma^2}(x-\mu)^2} dx = \frac{1}{\sqrt{2\pi}\sigma} \int e^{-\frac{1}{2\sigma^2}(x^2-2\mu x-2tx\sigma^2+\mu^2)} dx \\
 &= \frac{1}{\sqrt{2\pi}\sigma} \int e^{-\frac{1}{2\sigma^2}[\{x-(\mu+t\sigma^2)\}^2-(\mu+t\sigma^2)^2+\mu^2]} dx \\
 &= e^{\frac{1}{2\sigma^2}\{(\mu+t\sigma^2)^2-\mu^2\}} \frac{1}{\sqrt{2\pi}\sigma} \int e^{-\frac{1}{2\sigma^2}\{x-(\mu+t\sigma^2)\}^2} dx \\
 &= e^{\frac{1}{2\sigma^2}\{(\mu+t\sigma^2)^2-\mu^2\}} = e^{\mu t + t^2 \sigma^2 / 2}
 \end{aligned}$$

(iii) Standard normal distribution: $N(0, 1)$.

If $X \sim N(\mu, \sigma^2)$, $Z \equiv (X - \mu)/\sigma \sim N(0, 1)$. Hence

$$P(a < X < b) = P\left(\frac{a - \mu}{\sigma} < Z < \frac{b - \mu}{\sigma}\right) = \Phi\left(\frac{b - \mu}{\sigma}\right) - \Phi\left(\frac{a - \mu}{\sigma}\right),$$

where

$$\Phi(x) = P(Z < x) = \frac{1}{\sqrt{2\pi}} \int_{-\infty}^x e^{-u^2/2} du$$

is the CDF of $N(0, 1)$. Its values are tabulated in all statistical tables.

Example 3. Let $X \sim N(3, 5)$.

$$P(X > 1) = 1 - P(X < 1) = 1 - P\left(Z < \frac{1-3}{\sqrt{5}}\right) = 1 - \Phi(-0.8944) = 0.81.$$

Now find $x = \Phi^{-1}(0.2)$, i.e. x satisfies the equation


$$0.2 = P(X < x) = P\left(Z < \frac{x-3}{\sqrt{5}}\right).$$

From the normal table, $\Phi(-0.8416) = 0.2$. Therefore $(x-3)/\sqrt{5} = -0.8416$, leading to the solution $x = 1.1181$.


Note. You may check the answers using R:

```
1 - pnorm(1, 3, sqrt(5)),  
  qnorm(0.2, 3, sqrt(5))
```

Density functions of $N(0, \sigma^2)$



Density functions of $N(\mu, 1)$



Below are the R codes which produce the two normal density plots.

```
x <- seq(-4, 4, 0.1)      # x = (-4, -3.9, -3.8, ..., 3.9, 4)
plot(x, dnorm(x, 0, 1), type='l', xlab='x', ylab='Normal density',
      ylim=c(0, 0.85))
text(0,0.43, expression(sigma==1))
lines(x, dnorm(x, 0, 2), lty=2)
text(0,0.23, expression(sigma==sqrt(2)))
lines(x, dnorm(x, 0, 0.5), lty=4)
text(0,0.83, expression(sigma==sqrt(0.5)))
```

```
x <- seq(-3, 3, 0.1)
plot(x, dnorm(x, 0, 1), type='l', xlab='x', ylab='Normal density',
      xlim=c(-5, 5))
text(0,0.34, expression(mu==0))
lines(x+2, dnorm(x+2, 2, 1), lty=2)
text(2,0.34, expression(mu==2))
lines(x-2, dnorm(x-2, -2, 1), lty=4)
text(-2,0.34, expression(mu== -2))
```

Exponential Distribution $\text{Exp}(\lambda)$: $X \sim \text{Exp}(\lambda)$ if X has the PDF

$$f(x) = \begin{cases} \frac{1}{\lambda} e^{-x/\lambda} & x > 0 \\ 0 & \text{otherwise,} \end{cases}$$

where $\lambda > 0$ is a parameter.

$$E(X) = \lambda, \quad \text{Var}(X) = \lambda^2, \quad \psi_X(t) = 1/(1 - t\lambda).$$

Background. $\text{Exp}(\lambda)$ is used to model the lifetime of electronic components and the waiting times between rare events.

Gamma Distribution $\text{Gamma}(\alpha, \beta)$: $X \sim \text{Gamma}(\alpha, \beta)$ if X has the PDF

$$f(x) = \begin{cases} \frac{1}{\beta^\alpha \Gamma(\alpha)} x^{\alpha-1} e^{-x/\beta} & x > 0 \\ 0 & \text{otherwise,} \end{cases}$$

where $\alpha, \beta > 0$ are two parameters, and $\Gamma(\alpha) = \int_0^\infty y^{\alpha-1} e^{-y} dy$.

$$E(X) = \alpha\beta, \quad \text{Var}(X) = \alpha\beta^2, \quad \psi_X(t) = (1 - t\beta)^{-\alpha}.$$

Note. Gamma(1, β) = Exp(β).

Cauchy Distribution: the PDF of the Cauchy distribution is

$$f(x) = \frac{1}{\pi(1 + x^2)}, \quad x \in (-\infty, \infty).$$

As $E(|X|) = \infty$, the mean and variance of the Cauchy distribution do not exist. Cauchy Distribution is particularly useful to model the data with excessively large, or negatively large outliers.